



Visualizing Deep Learning using PyTorch and PowerAI

Dustin VanStee
Data Scientists
IBM Worldwide Client Experience Centers

2019 Yale University
Dec 3 | New Haven, CT



Agenda


- **Pytorch on PowerAI**
 - **What is Pytorch?**
 - **Where does it fit in deep learning frameworks**
 - **Where can I get it?**
- **Some examples**
- **GPU access**
- **Demo**
- **Lab !**

Deep Learning Frameworks

The word "Caffe" in a dark red, serif font.The PyTorch logo, featuring the word "PYTORCH" in a bold, black, sans-serif font, with a stylized orange flame icon replacing the letter "O".The DL4J logo, featuring the text "DL4J" in a large, black, sans-serif font, with "DEEPLARNING4J" in a smaller font below it.

Deep Learning Framework Popularity

Deep Learning Framework Six-Month Growth Scores 2019



- <https://towardsdatascience.com/which-deep-learning-framework-is-growing-fastest-3f77f14aa318>

PYTORCH – a quick review

Facebook's framework for research

- ✓ Cousin of LUA based Torch framework, but was rewritten to be tailored to Python frontend
- ✓ Gaining popularity quickly for its ease of use in R&D
- ✓ Supports dynamic computation graphs !
- ✓ Based on Python with Numpy compatibility
- ✓ Multi-GPU
- ✓ Easy to use, and supports standard debug tools



PyTorch Build	Stable (1.0)			Preview (Nightly)				
Your OS	Linux		Mac		Windows			
Package	Conda		Pip		LibTorch		Source	
Language	Python 2.7		Python 3.5		Python 3.6		Python 3.7	C++
CUDA	8.0		9.0		10.0		None	
Run this Command:	conda install pytorch torchvision cudatoolkit=10.0 -c pytorch							

PyTorch Basics – ND arrays


ND Tensors

When working with neural networks, you are always dealing with multidimensional arrays. Here are some quick tricks

Assume A is a 32x32 RGB image

```
## 3D Tensors
import torch
A = torch.rand(32,32,3)
plt.imshow(A)
```

<matplotlib.image.AxesImage at 0x10007f7ac208>



PyTorch Basics – Basic Slicing /Dimension Operations

Slicing Tensors - grab 'RED' dimension

```
In [15]: red_data = A[:, :, 0] #0 represents the first channel of RGB  
red_data.size()
```

```
Out[15]: torch.Size([32, 32])
```

Swap the RGB dimension and make the tensor a 3x32x32 tensor

```
In [11]: A_rgb_first = A.permute(2, 0, 1)  
print(A_rgb_first.size())
```

```
torch.Size([3, 32, 32])
```

Add a BatchSize to our Image Tensor

Usually you need to do this to run inference on your trained model

```
In [20]: Anew = A.unsqueeze(0)  
print(Anew.size())
```

```
torch.Size([1, 32, 32, 3])  
torch.Size([32, 32, 3])
```

PyTorch Basics – Basic Matrix Operations

Matrix Multiply

```
In [18]: a = torch.randint(4,(2,3))  
b = torch.randint(4,(3,2))  
print(a)  
print(b)
```

```
tensor([[0, 0, 1],  
        [0, 2, 1]])  
tensor([[1, 1],  
        [3, 3],  
        [0, 1]])
```

```
In [19]: # 2x3 @ 3x2 ~ 2x2  
a.matmul(b)  
torch.matmul(a,b)
```

```
Out[19]: tensor([[0, 1],  
                 [6, 7]])
```


PyTorch Basics – Basic Index Operations

Create a onehot vector

In [40]:

```
batch_size = 5
nb_digits = 10
# Dummy input that HAS to be 2D for the scatter (you can use view(-1,1) if needed)
y = torch.LongTensor(batch_size,1).random_() % nb_digits
# One hot encoding buffer that you create out of the loop and just keep reusing
y_onehot = torch.FloatTensor(batch_size, nb_digits)

# In your for loop
y_onehot.zero_()
y_onehot.scatter_(1, y, 1)

print(y)
print(y_onehot)
```

```
tensor([[ 8],
        [ 1],
        [ 4],
        [ 5],
        [ 7]])
tensor([[ 0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.],
        [ 0.,  1.,  0.,  0.,  0.,  0.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.,  0.,  0.],
        [ 0.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  0.,  0.]])
```

PyTorch: Using GPU's

CUDA Tensors

Tensors can be moved onto any device using the `.to` method.

```
In [17]: # let us run this cell only if CUDA is available
# We will use ``torch.device`` objects to move tensors in and out of GPU
x = torch.rand(2,2,2)
if torch.cuda.is_available():
    device = torch.device("cuda")           # a CUDA device object
    y = torch.ones_like(x, device=device)   # directly create a tensor on GPU
    x = x.to(device)                        # or just use strings ``.to("cuda")``
    z = x + y
    print(z)
    print(z.to("cpu", torch.double))       # ``.to`` can also change dtype together!

tensor([[[[1.5508, 1.0580],
          [1.6207, 1.1363]],

         [[1.9148, 1.6978],
          [1.5459, 1.5224]]], device='cuda:0'])
tensor([[[[1.5508, 1.0580],
          [1.6207, 1.1363]],

         [[1.9148, 1.6978],
          [1.5459, 1.5224]]], dtype=torch.float64])
```

PyTorch: Build a model

To build a neural network using pytorch, you need to perform three main steps

- extend nn.Module
- define `init` function
- define forward function

```
In [ ]: class NetCNN3L(nn.Module):

    def __init__(self):
        super(NetCNN3L, self).__init__()
        self.name = "NetCNN3L"
        # output dimension (H,W) -> H - kernel + 1 - 2p
        # in_NCHW=[Nx3x32x32 image], 3x3 square kernel, out_NCHW=[Nx32x32x32 image]
        # cin=3, cout=32
        self.conv1_1 = nn.Conv2d(3, 32, kernel_size=(3,3),padding=(1,1)) # same padding
        # in_NCHW=[Nx3x32x32 image], 3x3 square kernel, out_NCHW=[Nx32x32x32 image]
        self.conv1_2 = nn.Conv2d(32, 32, kernel_size=(3,3),padding=(1,1)) # same padding
        # in_NCHW=[Nx32x16x16 image], 3x3 square kernel, out_NCHW=[Nx64x16x16 image]
        self.conv2_1 = nn.Conv2d(32, 64, kernel_size=(3,3),padding=(1,1))
        # in_NCHW=[Nx64x16x16 image], 3x3 square kernel, out_NCHW=[Nx64x16x16 image]
        self.conv2_2 = nn.Conv2d(64, 64, kernel_size=(3,3),padding=(1,1))

        # an affine operation: y = Wx + b
        # 64 x 8 x 8
        self.fc1 = nn.Linear(4096, 512)
        self.fc2 = nn.Linear(512, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = self.conv1_1(x); self.c1_1 = x # (for plotting)
        x = F.relu(x)
        x = self.conv1_2(x); self.c1_2 = x # (for plotting)
        x = F.relu(x)
        x = F.max_pool2d(x, (2,2))
        x = F.dropout(x, p=0.25)
        x = self.conv2_1(x); self.c2_1 = x # (for plotting)
        x = F.relu(x)
        x = self.conv2_2(x); self.c2_2 = x # (for plotting)
        x = F.relu(x)
        x = F.max_pool2d(x, (2,2))
        x = F.dropout(x, p=0.25)
        #Flatten x for fully connected layer
        x = x.view(-1, 4096)
        #print(x.size())
        x = self.fc1(x)
        #print(x.size())
        x = F.relu(x)
        x = F.dropout(x, p=0.5)
        x = self.fc2(x)
        x = F.softmax(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

PyTorch: Train a model

```
for i,(X,Y) in enumerate(data_loader):
```

The key step during training is that after running an inference on a batch

```
yhat = model(X)
```

We calculate the loss

```
loss = criterion(yhat, Y)
```

and then update the weights based on the calculated gradients

```
loss.backward()
```

```
optimizer.step()
```

All the rest of this function is just mainly used for book keeping ...

PyTorch: Inference on a model

```
In [22]: # Send x to the GPU and run inference
# Select an image from our test set
IMG = 17
x = torch.tensor(np.asarray(test_set[IMG][0]), dtype=torch.float32)

# Plot the image (requires H,W,C as input so need to permute the axes)
plot_image(x.permute(1,2,0))


# For inference, convert C,H,W to N,C,H,W by adding the batch dimension . Batch size=1
x = x.unsqueeze(0) # add an batch dimension

y = model(x.to('cuda'))
y

torch.Size([32, 32, 3])
```

```
/gpfs/gpfs_gl4_16mb/s4s004/vansteer/anaconda3/envs/powerai-1.6.0/lib/python3.6/site-packages/invoke_launcher.py:59:
UserWarning: Implicit dimension choice for softmax has been deprecated.
nt.
```

```
Out[22]: tensor([[0., 0., 0., 0., 0., 0., 0., 1., 0., 0.], device='cuda:0',
grad_fn=<SoftmaxBackward>)
```



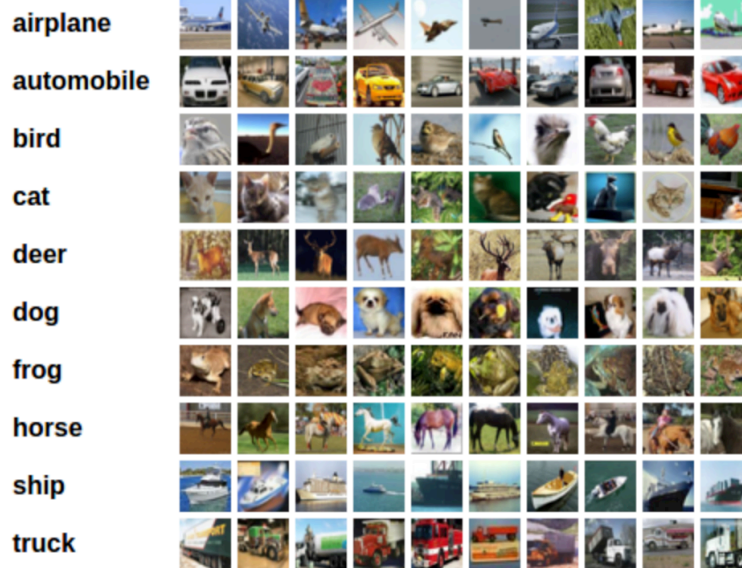
```
In [6]: labels_map = {
    0: "plane",
    1: "car",
    2: "bird",
    3: "cat",
    4: "deer",
    5: "dog",
    6: "frog",
    7: "horse",
    8: "ship",
    9: "truck"
}
```

PyTorch : TorchVision Library

Torchvision : An image library that contains

- Popular Image Datasets
- Popular NN Models (pretrained)
- Transformers

- AlexNet
- VGG
- ResNet
- SqueezeNet
- DenseNet
- Inception v3



CIFAR 10 Dataset

PyTorch : Vision Dataset API Example

```
xt = transforms.Compose([
    #transforms.CenterCrop(10),
    transforms.ToTensor(),
])
```

Torchvision : An image library that contains

- Popular Image Datasets
- Popular NN Models (pretrained)
- Transformers

```
train_fm_set = dset.FashionMNIST(root = './fashion_mnist_data',
                                  download=True, transform=xt)
test_fm_set = dset.FashionMNIST(root='./fashion_mnist_data', train=False,
                                 download=True, transform=xt) #transform=trans
```

```
|
test_fm_loader = torch.utils.data.DataLoader(
    dataset=test_fm_set,
    batch_size=batch_size,
    shuffle=False)
```



Pytorch PowerAI LAB –

Dustin VanStee
Bob Chesebrough
IBM WW Client Centers

Lab Code Here :

Code is here

<https://github.com/dustinvansteep/pytorch-examples>

